

# Curiously Recurring C++ Bugs

at Facebook

Louis Brandy  
Facebook

# C++ @ fb

- Second most used language at Facebook
- C++ in Infrastructure from very early on
- Now: C++, python, java, ...

# my team

- Responsible for the C++ tooling and core libraries
- Tools: compilers, sanitizers, CI, build, etc.
- Core libraries: folly, "common", refactors
- Frequently involved in nasty bugs, escalations.

# Agenda

- Bugs.
- Scary, common, concrete.
- Mitigation (welcome to the church of ASAN).
- **Warning:** gcc/clang bias from me.

# Two audiences

- Audience #1. Show you a bug you've never seen before.
  - Save you from a really, really bad day.
- Audience #2. Show you a bug you've seen a hundred times
  - that you never, ever make yourself anymore...

# Bug #1

```
std::vector::operator[]
```

# Bug #1

```
std::vector::operator[]  
    array[]  
    *ptr
```

# vector::operator[]

```
return vec[27];
```

- Crashing the Site 101.
- Eternal problem.
- ... and for our purposes, generally uninteresting.



# mitigation

- Static analysis. Complicated, hard, expensive.
- Improved abstractions (range-based operations).
- Dynamic analysis.
  - Bounds checking
  - Address sanitizer (+fuzzing!)
  - `-fsanitize=address`

# Bug #2

greatest C++ newbie bug

# Bug #2

```
std::map::operator[]
```

# std::map::operator[]

```
map<string, int> m;  
m["hey"] = 12;  
cout << m["hye"];
```

# Weird, but useful

```
map<char, int> occ;  
for (char c : str) {  
    occ[c]++;  
}
```

# Won't compile

## const correctness

```
void Widget::configure(  
    const map<string, string>& settings) {  
    m_timeout = settings["timeout"];  
    m_size = settings["max_items"];  
}
```

# Nooooooo!

```
void Widget::configure(  
    map<string, string>& settings) {  
    m_timeout = settings["timeout"];  
    m_size = settings["max_items"];  
}
```

# error message

```
error: passing 'const std::map<std::string<char>,  
                                std::string<char> >'  
as 'this' argument discards qualifiers
```



# "error passing const"

```
error: passing 'const std::map<std::string<char>,  
                                std::string<char> >'  
as 'this' argument discards qualifiers
```

# Twice

```
Widget::Widget(  
    const map<string, int>& settings) :  
    m_settings (settings) {  
  
  
  
  
  
  
  
  
  
}
```

# Twice

```
Widget::Widget(  
    const map<string, int>& settings) :  
    m_settings (settings) {  
    std::cout << "initializing with ... \n"  
               << "timeout: "  
               << m_settings["timeout"];  
               << std::endl;  
    }
```

# Helpfully set timeout to 0

```
Widget::Widget(  
    const map<string, string>& settings) :  
    m_settings (settings) {  
    std::cout << "initializing with ... \n"  
        << "timeout: "  
        << m_settings["timeout"];  
        << std::endl;  
    }
```

# mitigation

- none
- const correctness?
- ban it?



# Bug #3

`mapGetDefault()`

# some haskell!

```
findWithDefault def k m
  = case lookup k m of
      Nothing -> def
      Just x   -> x
```

# folly::get\_default

```
template <class Map, typename Key>
typename Map::mapped_type get_default(
    const Map& map,
    const Key& key,
    const typename Map::mapped_type& dflt ) {
    auto pos = map.find(key);
    return (pos != map.end() ? pos->second : dflt);
}
```



# string-version

```
string get_default(  
    const map<string,string>& map,  
    const string& key,  
    const string& dflt) {  
  
    auto pos = map.find(key);  
    return (pos != map.end() ?  
            pos->second : dflt);  
}
```

# fixed!

```
const string& get_default(  
    const map<string,string>& map,  
    const string& key,  
    const string& dflt) {  
  
    auto pos = map.find(key);  
    return (pos != map.end() ?  
            pos->second : dflt);  
}
```

## Sorry, this page isn't available

The link you followed may be broken, or the page may have been removed.



People love temporary  
default values.  
(understandably)

```
auto& v = get_default(m, k, "127.0.0.1");
```

# Broader class of problems

- "smuggling" a reference to a temporary through a function.
- One of many, many examples.
- C++ standards

# mitigation

- Address sanitizer will catch this!
- Maybe with extra flags. Depends on clang version. I think.
- `-fsanitize-address-use-after-scope`

# Bug #4

a success story

volatile



# volatile lint warning

**'volatile' does not make your code thread-safe. If multiple threads are sharing data, use `std::atomic` or locks. In addition, 'volatile' may force the compiler to generate worse code than it could otherwise.**

**we had so many arguments**

# wild success

- `std::atomics`
- higher level abstractions
- Solved a very hard, social problem.
- `expectations == reality?`

# Bug #5

the worst question in all of C++

Is

Is shared\_ptr

**Is `shared_ptr` thread-safe?**



is shared\_ptr|



is shared\_ptr **thread safe**

is\_shared\_ptr

is shared\_ptr **slow**

is shared\_ptr **null**

Google Search

I'm Feeling Lucky


*Report inappropriate predictions*




# Correct answer

- Is `shared_ptr` thread-safe?
- If you need to ask...
- Let's go with "No!"

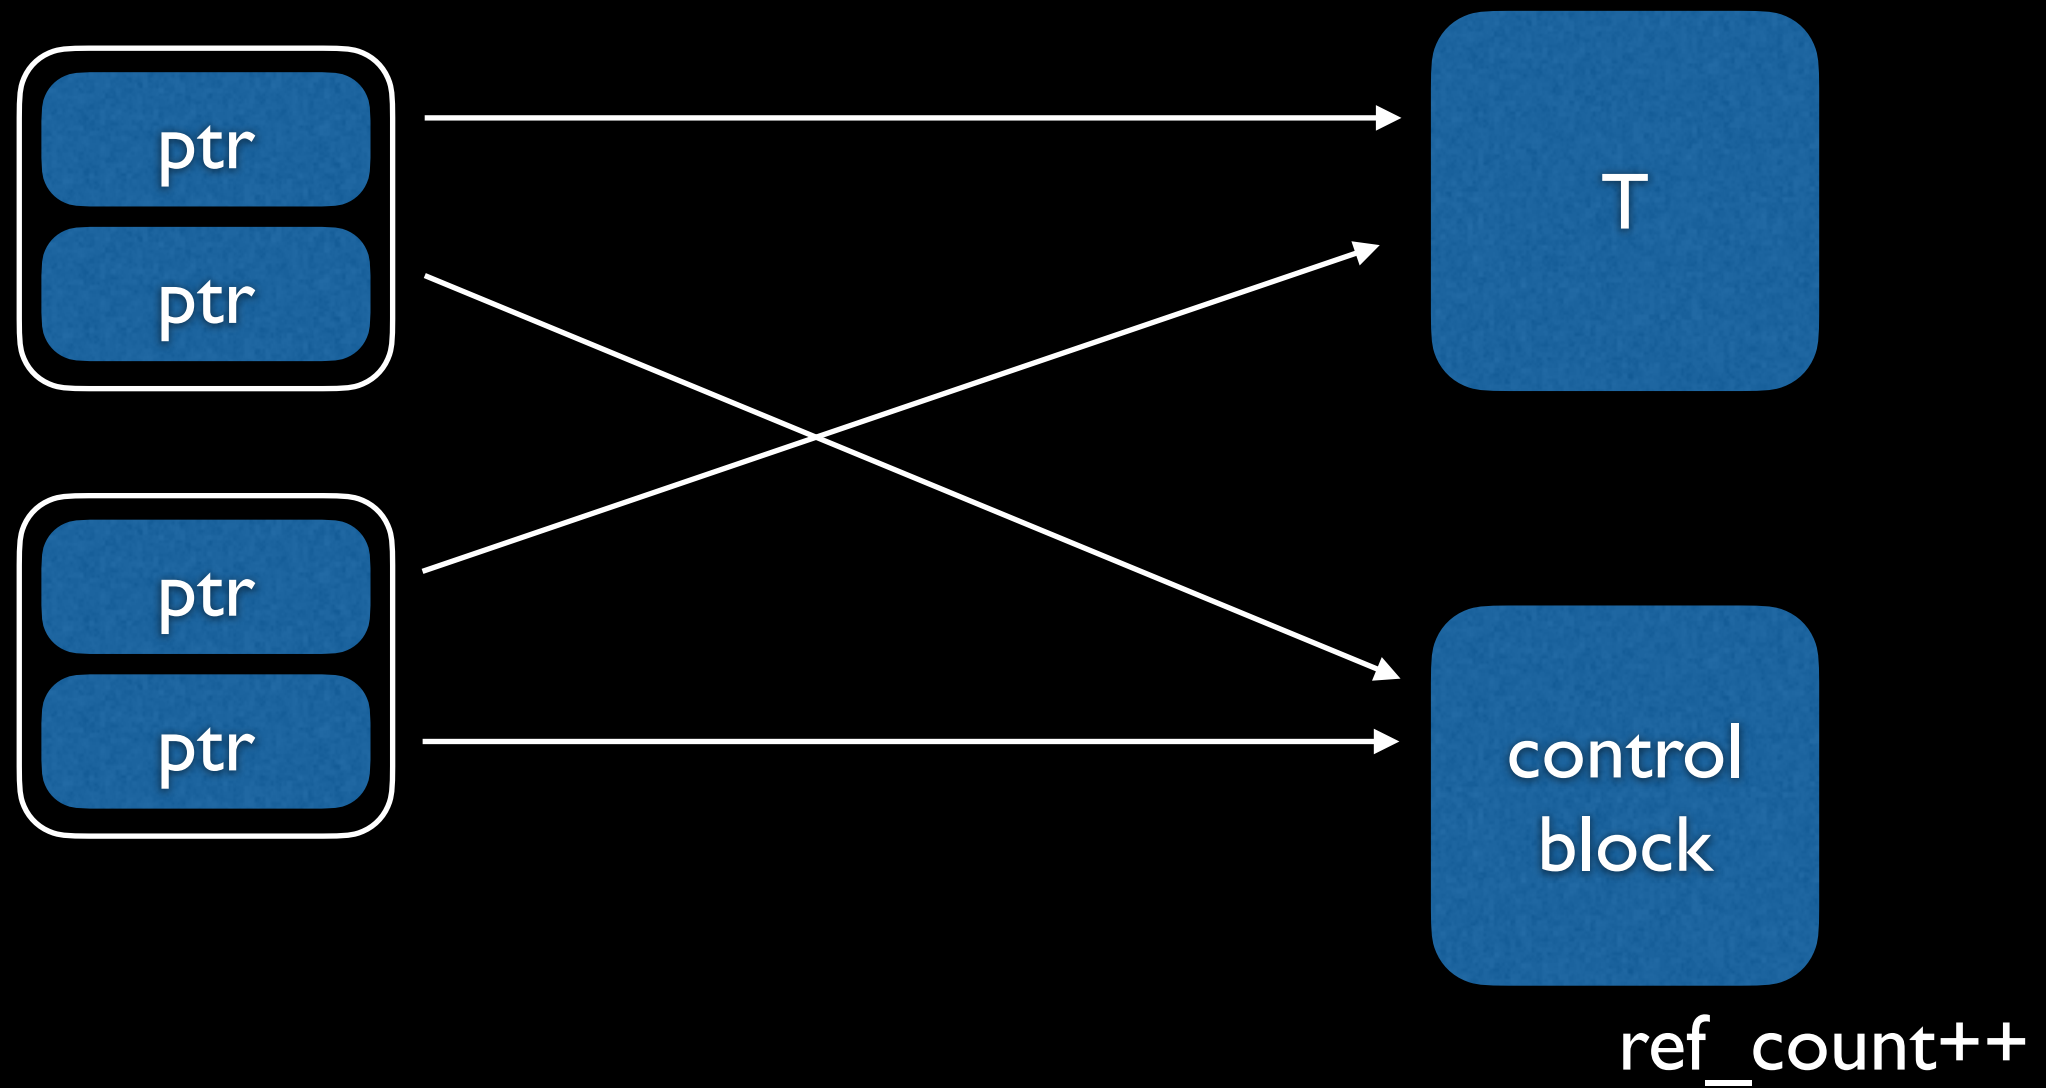
# a naive mental model




# thread safety





# copying




# thread safety




-  Thread safe
-  NOT thread safe

# Our intuition



Expected



Not so expected

# buggy RCU

- Read-Copy-Update
- One thread is writing (updating) to some one-true-shared\_ptr
- Many threads are reading (copying) the shared\_ptr.
- Use a lock on the shared\_ptr at the center.

# mitigation

- Thread sanitizer.
- Address Sanitizer often too, though incidentally.
- Library!
  - C++11: `atomic<shared_ptr>`
  - Concurrency TS: `atomic_shared_ptr`
  - Maybe back? `atomic<shared_ptr>`



# bonus bug

```
auto& ref = *returns_a_shared_ptr();  
ref.boom();
```

Sanitizers will catch this too.

# Bug #6

code review

# does this compile?

```
#include <string>
```

```
void f() {  
    std::string(foo);  
}
```

# does this compile?

```
#include <string>
```

```
void f() {  
    std::string(foo);  
}
```

✓ YES

# Same.

```
std::string(foo);
```

```
std::string foo;
```

x86-64 gcc 7.1 (Editor #1, Compiler #1) x

x86-64 gcc 7.1 -std=c++11

11010 .LX0: .text // \s+ Intel A

```
1 f():
2     push    rbp
3     mov     rbp, rsp
4     sub     rsp, 32
5     lea     rax, [rbp-32]
6     mov     rdi, rax
7     call    std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::basic_string()
8     lea     rax, [rbp-32]
9     mov     rdi, rax
10    call    std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::~~basic_string()
11    nop
12    leave
13    ret
```

# Ok, but so what?

```
std::string(foo);
```

```
std::string foo;
```

# See the problem?

```
void Obj::update() noexcept {
```



# See the problem?

```
void Obj::update() noexcept {  
    unique_lock<mutex>(m_mutex);  
    do_the_mutation();  
}
```

# Fixed

```
void Obj::update() noexcept {  
    unique_lock<mutex> g(m_mutex);  
    do_the_mutation();  
}
```

**Do a code search.**

**You'll find em.**

**Do a code search.**

**While prepping these slides...**

# I found a bug.

```
std::unique_lock<std::mutex>(mtx);
```

```
126 std::chrono::steady_clock::time_point Reader::getLastDataTime() {  
127     std::unique_lock<std::mutex>(mtx);
```

**Louis Brandy (ldbrandy)** commented

*Line 127* [Previous](#) · [Next](#) · [Like](#) · [Reply](#)

just fyi this is a bug. It needs to have a name.

```
std::unique_lock<std::mutex> **lock** (mtx);
```

# I found two bugs.

```
unique_lock<mutex>(m_);
```

```
11 shared_ptr<Client> ChronosJobBotIt::getClient(  
12     const std::string& clusterName) {  
13     unique_lock<mutex>(m_);
```

**Louis Brandy (ldbrandy) commented**

[Line 13](#) [Previous](#) · [Next](#) · [Like](#) · [Reply](#)

just fyi, this is a bug... the local variable needs a name....

```
unique_lock<mutex> lock(m_);
```

the worst C++ quiz  
does it compile?

# Remember?

```
#include <string>

void f() {
    std::string(foo);
}
```

✓ YES



# does this compile?

```
#include <string>

void f() {
    std::string(foo);
    std::string("wow");
}
```

✓ YES

# does this compile?

```
#include <string>

void f() {
    std::string(foo);
    std::string(foo);
}
```

**X NO**

# does this compile?

```
#include <string>
```

```
void f() {  
    std::string(foo);  
    std::string{foo};  
}
```

✓ YES

# does this compile?

```
#include <string>

void f() {
    std::string(foo);
    {std::string(foo);}
}
```

✓ YES

# Some important notes

```
unique_lock<mutex>(m_mutex);
```

- RAI-like types are especially affected because you won't otherwise "use" them.
- Only works because of the default constructor.
  - `std::unique_lock` is affected.
  - `std::lock_guard` is not.
- RAI types + default constructors -> **DANGER**.

# Mitigation

- This is a shadowing bug. `-Wshadow` finds this.
- `-Wshadow` tends to be noisy and isn't used often.
- `-Wshadow-compatible` and `-Wshadow-compatible-local` do **not** find these.

```
<source>: In function 'void f()':  
10 : <source>:10:41: warning: declaration of 'm_mutex' shadows a global declaration [-Wshadow]  
    std::unique_lock<std::mutex>(m_mutex);  
                                ^  
5 : <source>:5:12: note: shadowed declaration is here  
    std::mutex m_mutex;  
    ~~~~~  
Compiler exited with result code 0
```

# Easily detectable

- ... but not widely detected.
- Our clang based linter finds these, now.
- Warn on extraneous parenthesis in declarations?
- Perhaps a `[[nodiscard]]` attribute for constructors?

# Some conclusions

- C++ is hard.
- Tools are our best teaching weapons.
- Invest in a good, extendable linter (e.g. clang-tidy)
- ASAN is life.



# Questions?